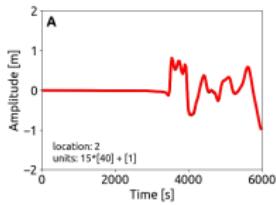
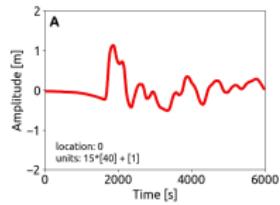
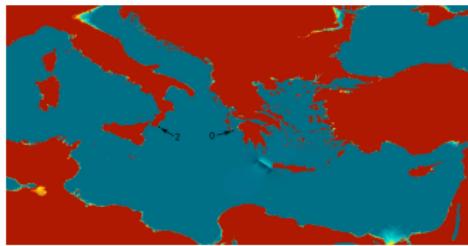


AI for Data-driven Simulations in Physics.

Siddhartha Mishra

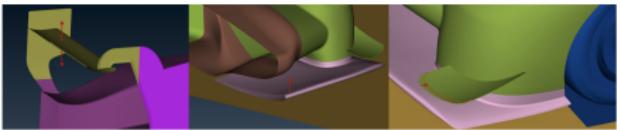
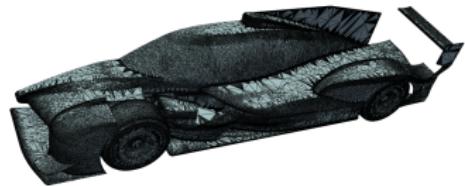
Computational and Applied Mathematics Laboratory (CamLab)
Seminar for Applied Mathematics (SAM), D-MATH (and),
ETH AI Center,
ETH Zürich, Switzerland.

Use Case I: Tsunami Early Warning System@INGV



- ▶ **Task:** Predict **Wave Height Time Series** at different Buoy locations in **Real Time**
- ▶ Basis of Tsunami Evacuation Forecast.

Use Case II: Race Car Design@Dallara

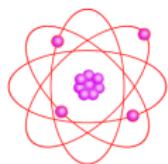


- ▶ Optimize Car Design.
- ▶ Predict **Aerodynamic body force** changes by changing specific parts.

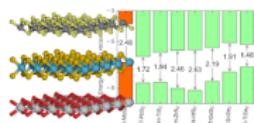
How are these problems solved currently ?

Step I: Mathematical Modeling

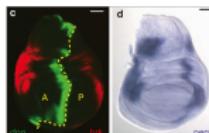
- ▶ Model Physical Phenomena with **Partial Differential Equations**
- ▶ PDEs are **Language of Nature**



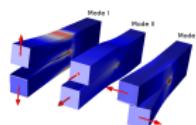
Schrödinger



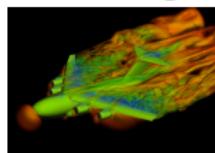
Kohn-Sham



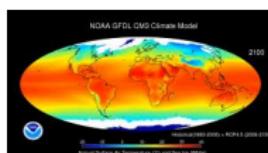
Reaction-Diffusion



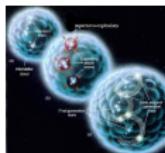
Phase-Field



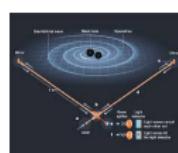
Euler



Navier-Stokes ++



MHD++

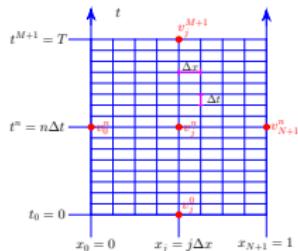


Einstein

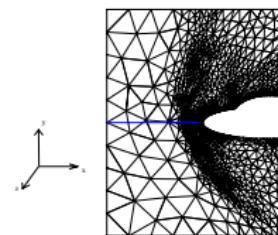
- ▶ Immense diversity of **Physical processes**
- ▶ Very wide range of **spatio-temporal scales**

Step II: Numerical Simulation

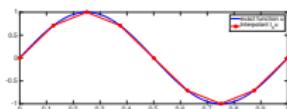
- ▶ Not possible to find solution formulas for PDEs.
- ▶ Reliance on **Numerical Methods** to approximate PDEs on computers.



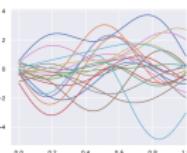
Finite Difference



Finite Volume



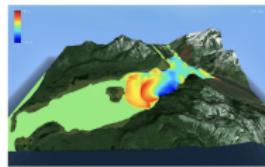
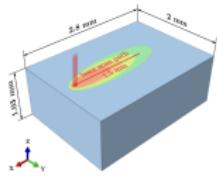
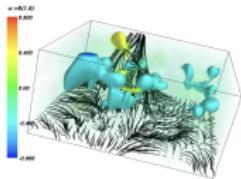
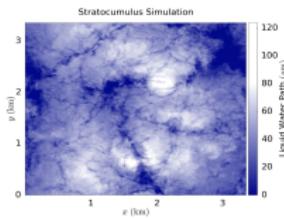
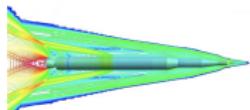
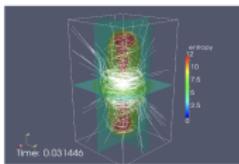
Finite Element



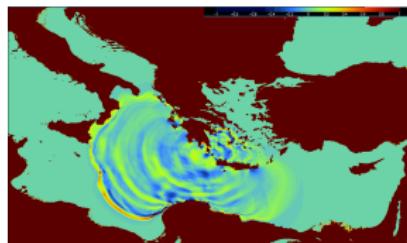
Spectral Method

Numerical Methods are very Successful

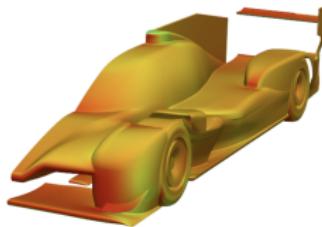
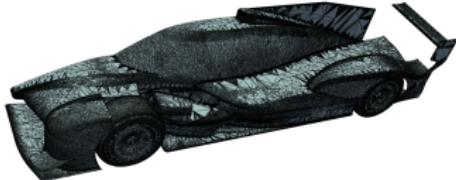
- ▶ Including @CAMIlab



What about the Use Cases ?

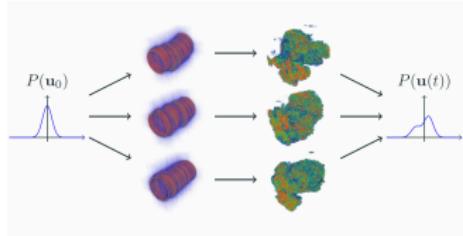
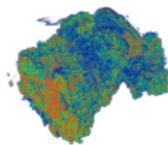
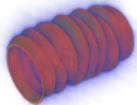


- ▶ Tsunami Simulation with Shallow-Water Equations
- ▶ Flow past Race car simulation with Navier-Stokes Equations



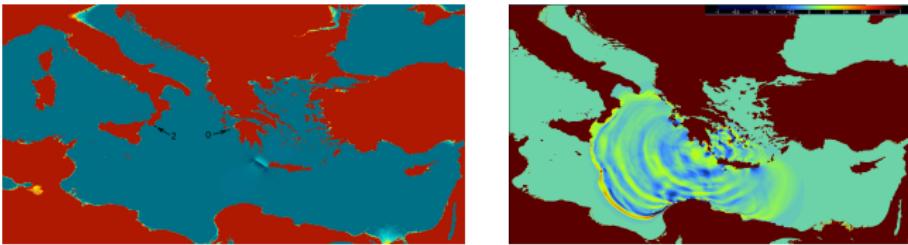
Where is the Catch ?

Issues with Numerical Methods I: Computational Cost



- ▶ PDE solvers can be very expensive,
- ▶ Many-Query Problems: UQ, Design, Inverse Problems.
- ▶ **Simulation** of Navier-Stokes at 1024^3 :
 - ▶ With **Azeban** on Piz Daint.
 - ▶ Single Run: 94 GPU hours (4512 CPU hours)
 - ▶ Ensemble simulation: 96256 node hours
 - ▶ Cost: Approx 500K USD.
 - ▶ **Solve PDEs fast**

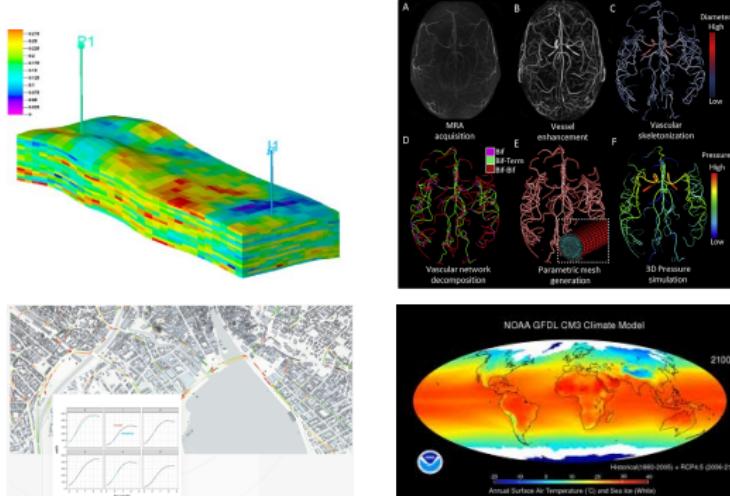
What about the Examples ?



- ▶ Single Tsunami Simulation takes > 1 hour !!
- ▶ Flow past Race car simulation requires 500 node hours per shape !!

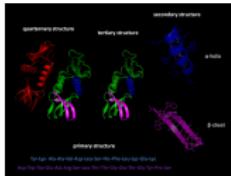


Issues with Numerical Methods II: Unknown Physics



- ▶ Missing Physics not just undetermined parameters.
- ▶ Manifestation of Sim2Real gap.
- ▶ Holds True for most real-world applications.
- ▶ Still have Data for the underlying Problem
- ▶ Learn PDE Solutions from Data + Physics

The age of AI

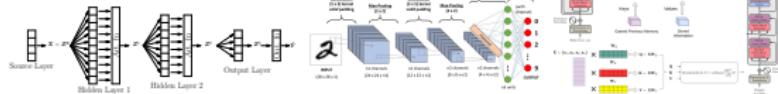


Machine learning for protein structure prediction
Deep learning models for protein structure prediction have revolutionized our ability to predict protein structures. This is a brief overview of how deep learning has changed the field.

- Traditional machine learning algorithms can accomplish a certain amount of protein structure prediction, but they are limited by the amount of training data available and the complexity of the underlying relationships between residues.
- Deep learning algorithms, on the other hand, can learn complex, non-linear relationships from large amounts of data, making them better suited for protein structure prediction.
- Protein structure prediction is a challenging task because proteins are composed of many different components, such as amino acids, which interact in complex ways. Deep learning models can learn these interactions directly from the data, without having to explicitly define them.
- Overall, deep learning models for protein structure prediction have revolutionized the field, making it easier and faster to predict protein structures. They have also opened up new applications, such as drug design and protein engineering, that were previously impossible or impractical.

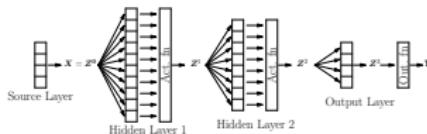


- ▶ Exponentially more **Compute** aka GPUs :-)
- ▶ Huge Data
- ▶ Deep Neural Networks



- Can Neural Networks solve PDEs ?

What are Deep Neural networks ?



- ▶ $\mathcal{L}^*(z) = \sigma_o \odot C_K \odot \sigma \odot C_{K-1} \dots \odot \sigma \odot C_2 \odot \sigma \odot C_1(z)$.
- ▶ At the k -th **Hidden layer**: $z^{k+1} := \sigma(C_k z^k) = \sigma(W_k z^k + B_k)$
- ▶ **Tuning Parameters**: $\theta = \{W_k, B_k\} \in \Theta$,
- ▶ σ : scalar **Activation function**: ReLU, Tanh
- ▶ **Random Training set**: $\mathcal{S} = \{z_i\}_{i=1}^N \in Z$, with i.i.d z_i
- ▶ Use **SGD** (ADAM) to find **Target** $\mathcal{L} \approx \mathcal{L}^* = \mathcal{L}_{\theta^*}^*$

$$\theta^* := \arg \min_{\theta \in \Theta} \sum_{i=1}^N |\mathcal{L}(z_i) - \mathcal{L}_\theta^*(z_i)|^p,$$

Physics Informed Neural Networks

- ▶ Variants of PINNs stem from [Dissanayake, Phan-Thien, 1994](#).
- ▶ Also in [Lagaris et al, mid 1990s](#).
- ▶ Reintroduced by [Raissi, Perdikaris, Karniadakis, 2017](#).
- ▶ Extensively developed by [Karniadakis](#) and collaborators.
- ▶ 10000s of papers on PINNs already.

PINNs for the PDE $\mathcal{D}(u) = f$

- ▶ For **Parameters** $\theta \in \Theta$, $u_\theta : \mathbb{D} \mapsto \mathbb{R}^m$ is a **DNN**, with $u_\theta \in X^*$
- ▶ Aim: Find $\theta \in \Theta$ such that $u_\theta \approx u$ (in suitable sense).
- ▶ Compute **PDE Residual** by Automatic Differentiation:

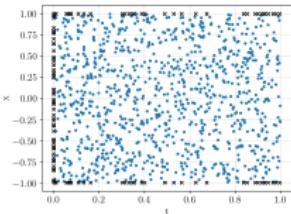
$$\mathcal{R} := \mathcal{R}_\theta(y) = \mathcal{D}(u_\theta(y)) - f(y), \quad y \in \mathbb{D} \quad \mathcal{R}_\theta \in Y^*, \quad \forall \theta \in \Theta$$

- ▶ **PINNs** are minimizers of $\|\mathcal{R}_\theta\|_Y^p \sim \int_{\mathbb{D}} |\mathcal{R}_\theta(y)|^p dy$
- ▶ Replace **Integral** by **Quadrature** !
- ▶ Let $\mathcal{S} = \{y_i\}_{1 \leq i \leq N}$ be quadrature points in \mathbb{D} , with weights w_i
- ▶ Could be Random, Sobol, Grid points (Gauss rules)
- ▶ **PINN** for approximating PDE is defined as $u^* = u_{\theta^*}$ such that

$$\theta^* = \arg \min_{\theta \in \Theta} \sum_{i=1}^N w_i |\mathcal{R}_\theta(y_i)|^p$$

Heat Eqn: $u_t = u_{xx}$ with 0-BC and $u(x, 0) = \bar{u}(x)$ IC

- ▶ Training Set: $\mathcal{S} = \mathcal{S}_{int} \cup \mathcal{S}_{tb} \cup \mathcal{S}_{sb}$ Randomly chosen.



- ▶ Deep Neural networks : $(x, t) \mapsto u_\theta(x, t)$, $\theta \in \Theta$.
- ▶ Temporal boundary residual: $\mathcal{R}_{tb,\theta} = u_\theta(\cdot, 0) - \bar{u}$
- ▶ Spatial boundary residual: $\mathcal{R}_{sb,\theta} = u_\theta|_{\partial D}$.
- ▶ Interior PDE Residual: $\mathcal{R}_{int,\theta} = \partial_t u_\theta - \partial_{xx} u_\theta$
- ▶ Evaluate PDE Residual by Automatic Differentiation
- ▶ Loss function:

$$J = \frac{1}{N_{tb}} \sum_{n=1}^{N_{tb}} |\mathcal{R}_{tb,\theta}(x_n)|^2 + \frac{1}{N_{sb}} \sum_{n=1}^{N_{sb}} |\mathcal{R}_{sb,\theta}(x_n, t_n)|^2 + \frac{1}{N_{int}} \sum_{n=1}^{N_{int}} |\mathcal{R}_{int,\theta}|^2.$$

Why PINNs are great ?: I

- ▶ Very easy to Code !!
- ▶ A few lines in PyTorch

```
def compute_res(self, network, x_f_train):
    x_f_train.requires_grad = True
    u = network(x_f_train).reshape(-1, )
    grad_u = torch.autograd.grad(u, x_f_train, grad_outputs=torch.ones(x_f_train.shape[0], ).to(self.device), create_graph=True)[0]
    grad_u_t = grad_u[:, 0]
    grad_u_x = grad_u[:, 1]
    grad_u_xx = torch.autograd.grad(grad_u_x, x_f_train, grad_outputs=torch.ones(x_f_train.shape[0], ).to(self.device), create_graph=True)[0][:, 1]

    residual = grad_u_t - self.v * grad_u_xx
    return residual
```

- ▶ Don't need Grids !!

Numerical Results: (SM, Molinaro, Tanios, 2021)

► Heat Equation:

Dimension	Training Error	Total error
20	0.006	0.79%
50	0.006	1.5%
100	0.004	2.6%

► Black-Scholes type PDE with Uncorrelated Noise:

Dimension	Training Error	Total error
20	0.0016	1.0%
50	0.0031	1.5%
100	0.0031	1.8%

► Heston option-pricing PDE

Dimension	Training Error	Total error
20	0.0064	1.0%
50	0.0037	1.3%
100	0.0032	1.4%

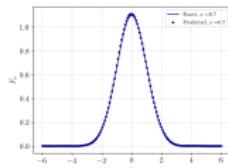
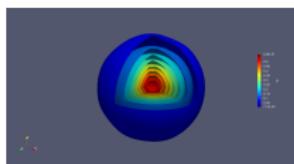
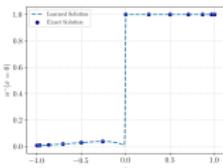
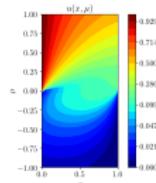
Radiative Transfer Equations

- ▶ 2d + 1-dim Integro-Differential PDE for Intensity

$$\frac{1}{c} u_t + \omega \cdot \nabla u + (k(x, \nu) + \sigma(x, \nu)) u - \frac{\sigma(x, \nu)}{s_d} \int_{R_+} \int_S \Phi(\omega, \omega', \nu, \nu') u d\omega' d\nu' = f(x, t, n, \nu).$$

- ▶ High-dimensional, non-local, mixed-type, multiphysics
- ▶ PINNs applied and bound derived in SM, Molinaro 2021.

Numerical Results



2-D, Intensity

2-D, Boundary

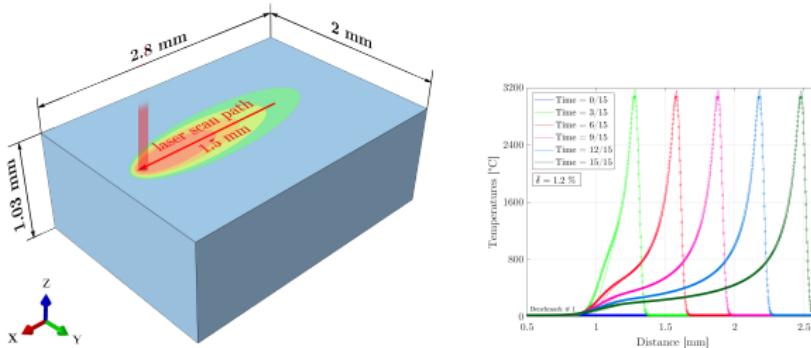
6-D, Inc. Radiation

6-D, Radial flux

Dimension	Network Size	Error	Training Time
2	24×8	0.3%	57 min
6	20×8	2.1%	66 min

An Industrial case study

- ▶ PINN simulation of Laser Powder Bed Fusion
- ▶ Key Component of 3-D Printing



- ▶ Traditional FEM Simulation: 4 hrs.
- ▶ PINNs of Hosseini et al, 2022: 2×10^{-6} secs with 4% Error.

Why do PINNs work or do they ?

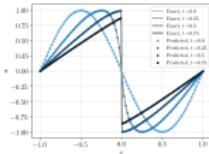
- ▶ Based on sound theory.
- ▶ Error Bounds of SM, Molinaro, De Ryck et.al 2021-2024:
- ▶ For generic PDE: $\mathcal{D}(u) = f$:

$$\|u - u_\theta\| \sim C_{\text{pde}}(u, u_\theta) [\mathcal{E}_T(\theta) + C_{\text{quad}}(u_\theta) N^{-\alpha}]$$

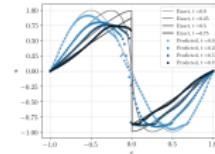
- ▶ C_{pde} depends on $\|\nabla u\|$.
- ▶ Can blow up for large gradients.

Viscous Burgers': $u_t + \operatorname{div} f(u) = \nu \Delta u$

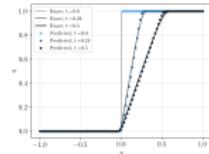
- ▶ Error $\mathcal{E} \leq C e^{CT} (\mathcal{E}_T + C_q N^{-\alpha})$, $C = C(\|\nabla u^\nu\|_{L^\infty})$
- ▶ $\|\nabla u^\nu\|_{L^\infty} \sim \frac{1}{\sqrt{\nu}} \Rightarrow$ Error can blow up near shocks !!



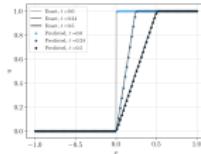
$\nu = 10^{-3}$, Sh



$\nu = 0$, Sh



$\nu = 10^{-3}$, RF



$\nu = 0$, RF

ν	\mathcal{E} (Shock)	\mathcal{E} (Rarefaction)
10^{-3}	1.0%	2.2%
10^{-4}	11.2%	1.6%
0	23.1%	1.2%

What about Training Error ?

- ▶ Rigorous Error estimate for PINNs for the PDE $\mathcal{D}(u) = f$:

$$\|u - u_\theta\| \sim C_{\text{pde}}(u, u_\theta) [\mathcal{E}_T(\theta) + C_{\text{quad}}(u_\theta) N^{-\alpha}]$$

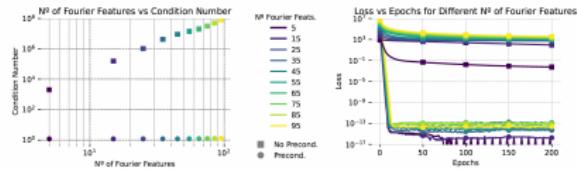
- ▶ Training Error is a **blackbox**
- ▶ We have that $\min_{\theta} \mathcal{E}_T(\theta) \leq \epsilon$
- ▶ But can we train to reach close to the global minimum with **Gradient Descent** ?
- ▶ De Ryck, SM et al (2024) showed that:

$$N(\delta) \sim \mathcal{O}(\kappa(\mathcal{A}) \log(1/\delta)), \quad \kappa(\mathcal{A}) = \frac{\lambda_{\max}(\mathcal{A})}{\lambda_{\min}(\mathcal{A})}, \quad \mathcal{A} \sim \mathcal{D}^* \mathcal{D}$$

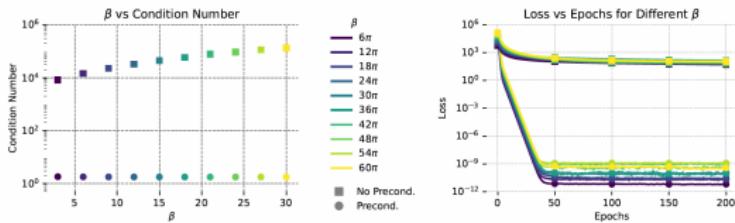
- ▶ Convergence of PINNs depends on **Conditioning of Hermitian-Square !!**
- ▶ Ex: if $\mathcal{D} = -\Delta$, then $\mathcal{A} = \Delta^2$

Training PINNs is ill-Conditioned

- For Poisson Equation: $-u'' = -k^2 \sin(kx)$:

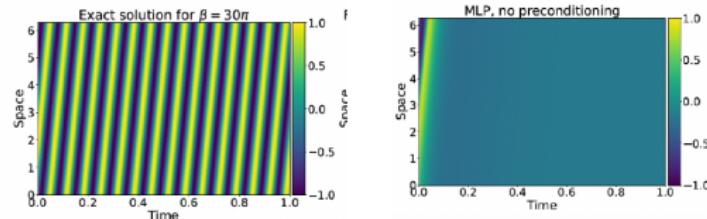


- For Advection Equation: $u_t + \beta u_x = 0$

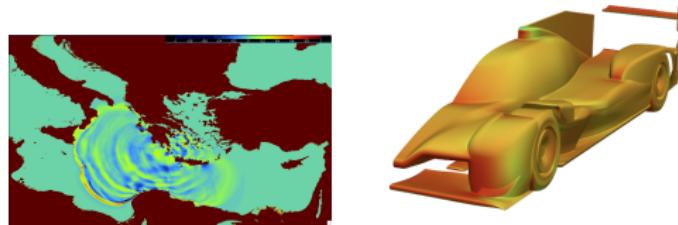


Intrinsic Limitations of PINNs

- ▶ Don't work on simple problems (Advection with $\beta = 30\pi$):



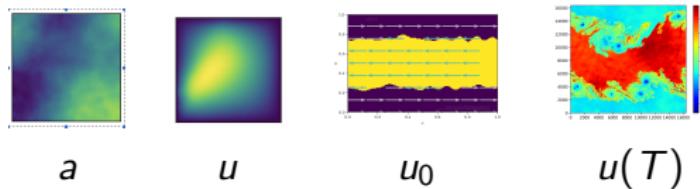
- ▶ Let alone real use cases !!



- ▶ Preconditioning is an active research area !!
- ▶ Have to bring Data to centerstage.

What does solving a PDE entail ?

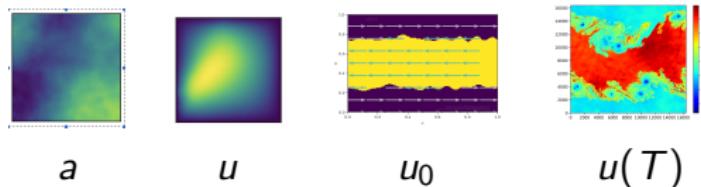
- ▶ Finding Solution Operators of PDEs.
- ▶ Darcy PDE: $-\operatorname{div}(a \nabla u) = f$, $\mathcal{G} : a \mapsto \mathcal{G}a = u$.



- ▶ Compressible Euler equations: $\mathcal{G} : u_0 \mapsto \mathcal{G}u_0 = u(t)$.
- ▶ Operator: $\mathcal{G} : \mathcal{X} \mapsto \mathcal{Y}$, $\dim(\mathcal{X}, \mathcal{Y}) = \infty$.
- ▶ Learn PDE Solution Operators from Data
- ▶ Underlying Data Distribution $\mu \in \operatorname{Prob}(\mathcal{X})$
- ▶ Draw N i.i.d samples $(a_i, \mathcal{G}(a_i))$ with $a_i \sim \mu$.
- ▶ Operator Learning Task: Find approximation to $\mathcal{G}_\# \mu$

What does solving a PDE entail ?

- ▶ Finding Solution Operators of PDEs.
- ▶ Darcy PDE: $-\operatorname{div}(a \nabla u) = f$, $\mathcal{G} : a \mapsto \mathcal{G}a = u$.



- ▶ Compressible Euler equations: $\mathcal{G} : u_0 \mapsto \mathcal{G}u_0 = u(t)$.
- ▶ Operator: $\mathcal{G} : \mathcal{X} \mapsto \mathcal{Y}$, $\dim(\mathcal{X}, \mathcal{Y}) = \infty$.
- ▶ Learn PDE Solution Operators from Data
- ▶ Underlying Data Distribution $\mu \in \text{Prob}(\mathcal{X})$
- ▶ Draw N i.i.d samples $(a_i, \mathcal{G}(a_i))$ with $a_i \sim \mu$.
- ▶ Operator Learning Task: Find approximation to $\mathcal{G}_\# \mu$
- ▶ Caveat: Neural Networks: $\mathbb{R}^N \mapsto \mathbb{R}^M$

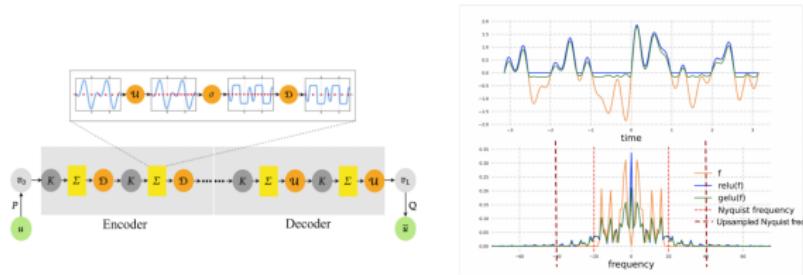
Possible Solution: Neural Operators

- ▶ DNNs are $\mathcal{L}_\theta = \sigma_K \odot \sigma_{K-1} \odot \dots \odot \sigma_1$
- ▶ Single hidden layer: $\sigma_k(y) = \sigma(A_k y + B_k)$, with $y \in \mathbb{R}^N$
- ▶ Generalize DNNs to ∞ -dimensions: Kovachki et al, 2021:
- ▶ NO: $\mathcal{N}_\theta = \mathcal{N}_L \odot \mathcal{N}_{L-1} \odot \dots \odot \mathcal{N}_1$
- ▶ Single hidden layer;

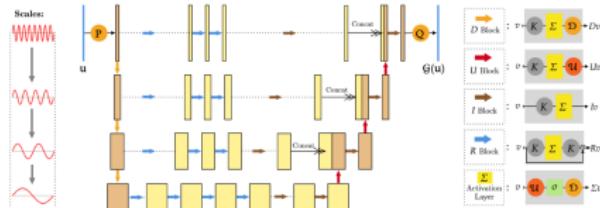
$$(\mathcal{N}_\ell v)(x) = [\mathcal{P}\sigma] \left(B_\ell(x) + \int_D K_\ell(x, y)v(y)dy \right)$$

- ▶ Kernel Integral Operators with Parameters B_ℓ, K_ℓ
- ▶ Nonlocal activations $[\mathcal{P}\sigma]$ Bartolucci, SM et. al, 2023

Convolutional Neural Operators: Raonic, SM et al, 2023

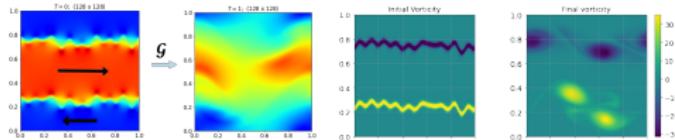


- ▶ I: Use **Convolutional Kernels in Physical space**
- ▶ II: Modulated Nonlocal activations for Alias-free processing.
- ▶ CNO instantiated as a modified **Operator UNet**

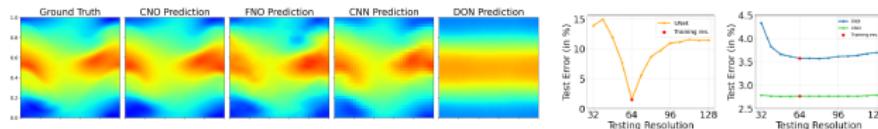


Example: Navier-Stokes Eqns.

- ▶ Operator:



- ▶ Comparison:



- ▶ Test Errors:

Model	FFNN	UNet	DeepONet	FNO	CNO
Error	8.05%	3.54%	11.64%	3.93%	3.01%

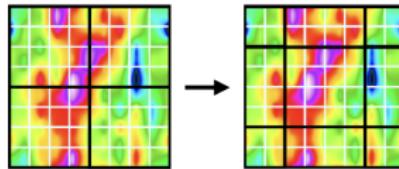
- ▶ CNO is Resolution Invariant a la Bartolucci, SM et. al, 2023

What about Nonlinear Kernels ?

- ▶ **Operator Attention:** $\mathbb{A}(v)(x) = \int_D K(v(x), v(y))v(y)dy :$

$$u(x) = \mathbb{A}(v)(x) = W \int_D \frac{e^{\frac{\langle Qv(x), Kv(y) \rangle}{\sqrt{m}}}}{\int_D e^{\frac{\langle Qv(z), Kv(y) \rangle}{\sqrt{m}}} dz} Vv(y) dy.$$

- ▶ Computational Cost is Quadratic in # (Tokens) !!
- ▶ Scaling through Vision + SWIN transformers



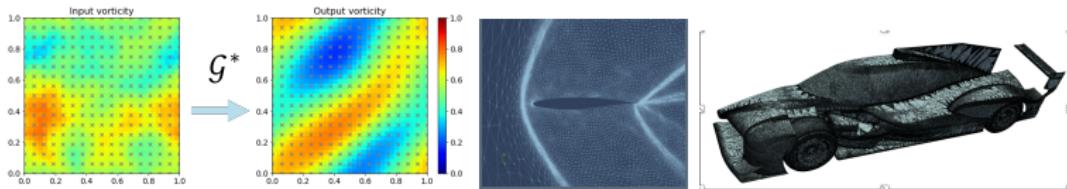
- ▶ scOT of Herde,SM et. al., 2024.

Models perform very well on 2-D Cartesian Domains !!

- ▶ Extensive Empirical evaluation on RPB benchmarks.

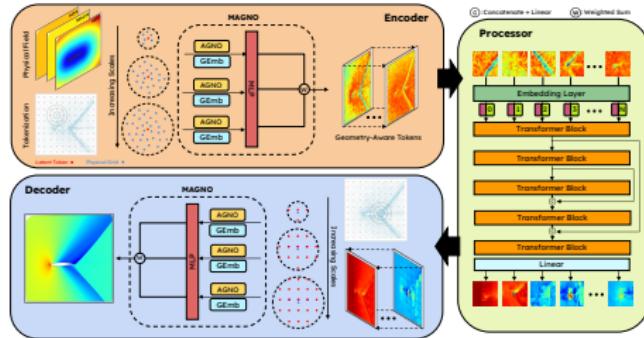
	In/Out	FFNN	GT	UNet	ResNet	DON	FNO	CNO
Poisson Equation	In	5.74%	2.77%	0.71%	0.43%	12.92%	4.98%	0.21%
	Out	5.35%	2.84%	1.27%	1.10%	9.15%	7.05%	0.27%
Wave Equation	In	2.51%	1.44%	1.51%	0.79%	2.26%	1.02%	0.63%
	Out	3.01%	1.79%	2.03%	1.36%	2.83%	1.77%	1.17%
Smooth Transport	In	7.09%	0.98%	0.49%	0.39%	1.14%	0.28%	0.24%
	Out	650.6%	875.4%	1.28%	0.96%	157.2%	3.90%	0.46%
Discontinuous Transport	In	13.0%	1.55%	1.31%	1.01%	5.78%	1.15%	1.01%
	Out	257.3%	22691.1%	1.35%	1.16%	117.1%	2.89%	1.09%
Allen-Cahn Equation	In	18.27%	0.77%	0.82%	1.40%	13.63%	0.28%	0.54%
	Out	46.93%	2.90%	2.18%	3.74%	19.86%	1.10%	2.23%
Navier-Stokes Equations	In	8.05%	4.14%	3.54%	3.69%	11.64%	3.57%	2.76%
	Out	16.12%	11.09%	10.93%	9.68%	15.05%	9.58%	7.04%
Darcy Flow	In	2.14%	0.86%	0.54%	0.42%	1.13%	0.80%	0.38%
	Out	2.23%	1.17%	0.64%	0.60%	1.61%	1.11%	0.50%
Compressible Euler	In	0.78%	2.09%	0.38%	1.70%	1.93%	0.44%	0.35%
	Out	1.34%	2.94%	0.76%	2.06%	2.88%	0.69%	0.59%

Caveat I: PDEs on Arbitrary Domains

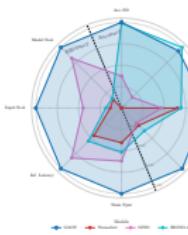
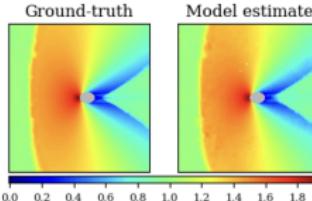
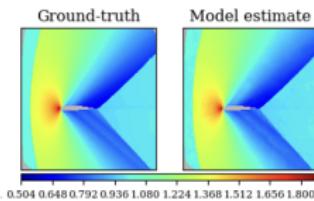


- ▶ Discussion so far has only focussed on **Cartesian Domains**
- ▶ Discretized with **Uniform Grids**.
- ▶ Most Real world PDEs are on **Arbitrary Domains**
- ▶ Discretized with **Unstructured Grids** or **Point Clouds**
- ▶ Need to handle such Data !!

Use Graphs + Transformers

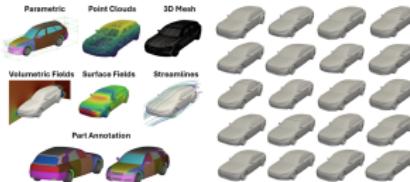


- ▶ **Geometry Aware Operator Transformer:** Wen, SM et. al, 2025
- ▶ GAOT is both accurate and efficient.

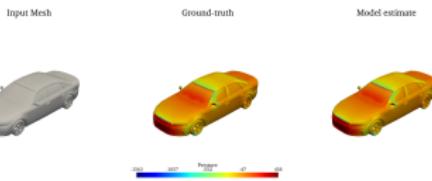


The DrivaerNet++ Challenge

- ▶ Flow past Cars Dataset (8K Car Shapes with 2M nodes each)



- ▶ GAOT: SOTA for Surface Pressure, Shear Stress

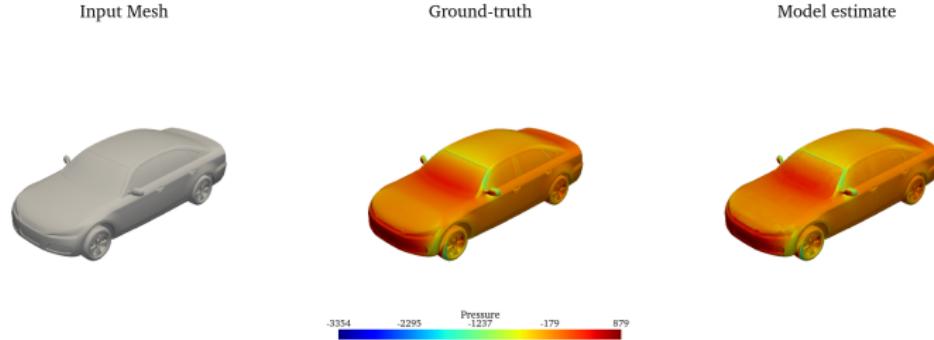


Model	GAOT	FigConvNet	TripNet	RegDGCNN
L^1 Pressure Err.	0.110	0.122	0.125	0.161
L^1 Shear Err.	0.156	0.222	0.215	0.364

- CFD: 300 Node hours vs. GAOT: 0.36 seconds !!!

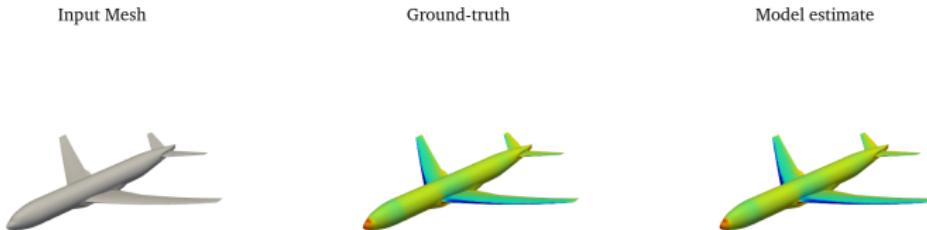
The DriverML challenge

- ▶ HR-LES simulations of flow past 500 cars.
- ▶ More accurate than RANS for Drivearnet++.
- ▶ Up to 10 M surface nodes handled accurately by GAOT !!



Flow Past an entire Aircraft

- ▶ AIAA's **NASA CRM** Benchmark.
- ▶ **GAOT** predicts Surface Pressure+Skin Friction accurately !!



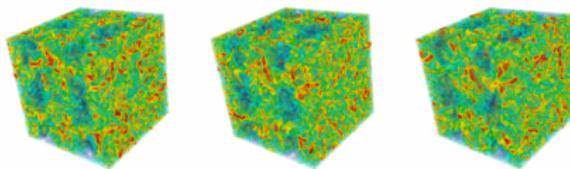
-1.80 -1.06 Pressure 0.40 1.33



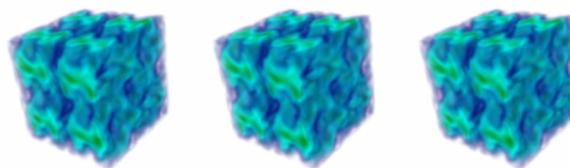
-2.56e-03 -2.22e-04 Surface Friction Coefficient (z-component) 2.10e-03 4.41e-03 6.73e-03

Caveat II: PDE with Chaotic Multiscale Solutions

- ▶ 3-D Navier-Stokes with Taylor-Green initial data.
- Spectral Viscosity Method:



- Convolutional Fourier Neural Operator (C-FNO):



- ▶ All ML models trained to minimize MSE or MAE:
 - ▶ Smooth out Small Scales
 - ▶ Collapse to Mean

- ▶ Insensitivity of Neural Networks:

$$\Psi_\theta(u + \delta u) \approx \Psi_\theta(u), \quad \delta u \ll 1$$

- ▶ DNNs are optimal at the Edge of Chaos: $\text{Lip}(\Psi_\theta) \sim \mathcal{O}(1)$
- ▶ Spectral Bias of DNNs
- ▶ Bounded Gradients are essential for training with GD
- ▶ Implication \Rightarrow DNNs will Collapse to Mean !!

$$\begin{aligned}\mathbb{E}_{\delta \bar{u}} \|\Psi_\theta(\bar{u}^* + \delta \bar{u}) - \mathcal{S}(\bar{u}^* + \delta \bar{u})\|^2 &\approx \mathbb{E}_{\delta \bar{u}} \|\Psi_\theta(\bar{u}^*) - \mathcal{S}(\bar{u}^* + \delta \bar{u})\|^2 \quad (\text{insensitivity hypothesis}) \\ &= \|\Psi_\theta(\bar{u}^*) - \mathbb{E}_{\delta \bar{u}} \mathcal{S}(\bar{u}^* + \delta \bar{u})\|^2 + \text{Var}_{\delta \bar{u}} [\mathcal{S}(\bar{u}^* + \delta \bar{u})]. \quad (\text{bias-variance decomposition})\end{aligned}$$

- ▶ Insensitivity of Neural Networks:

$$\Psi_\theta(u + \delta u) \approx \Psi_\theta(u), \quad \delta u \ll 1$$

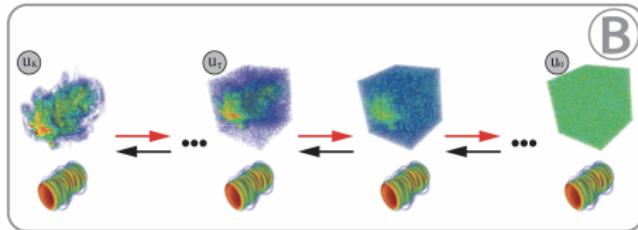
- ▶ DNNs are optimal at the Edge of Chaos: $\text{Lip}(\Psi_\theta) \sim \mathcal{O}(1)$
- ▶ Spectral Bias of DNNs
- ▶ Bounded Gradients are essential for training with GD
- ▶ Implication \Rightarrow DNNs will Collapse to Mean !!

$$\begin{aligned}\mathbb{E}_{\delta \bar{u}} \|\Psi_\theta(\bar{u}^* + \delta \bar{u}) - \mathcal{S}(\bar{u}^* + \delta \bar{u})\|^2 &\approx \mathbb{E}_{\delta \bar{u}} \|\Psi_\theta(\bar{u}^*) - \mathcal{S}(\bar{u}^* + \delta \bar{u})\|^2 \quad (\text{insensitivity hypothesis}) \\ &= \|\Psi_\theta(\bar{u}^*) - \mathbb{E}_{\delta \bar{u}} \mathcal{S}(\bar{u}^* + \delta \bar{u})\|^2 + \text{Var}_{\delta \bar{u}} [\mathcal{S}(\bar{u}^* + \delta \bar{u})]. \quad (\text{bias-variance decomposition})\end{aligned}$$

- Directly Learn the Conditional Distribution $\mathcal{S}_\# \mu$

GenCFD algorithm of Molinaro et. al, SM, 2025

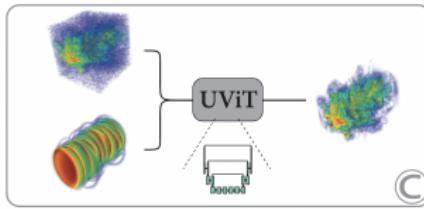
- Based on **Conditional Score Based Diffusion Models**



- Denoised with the **Reverse SDE**:

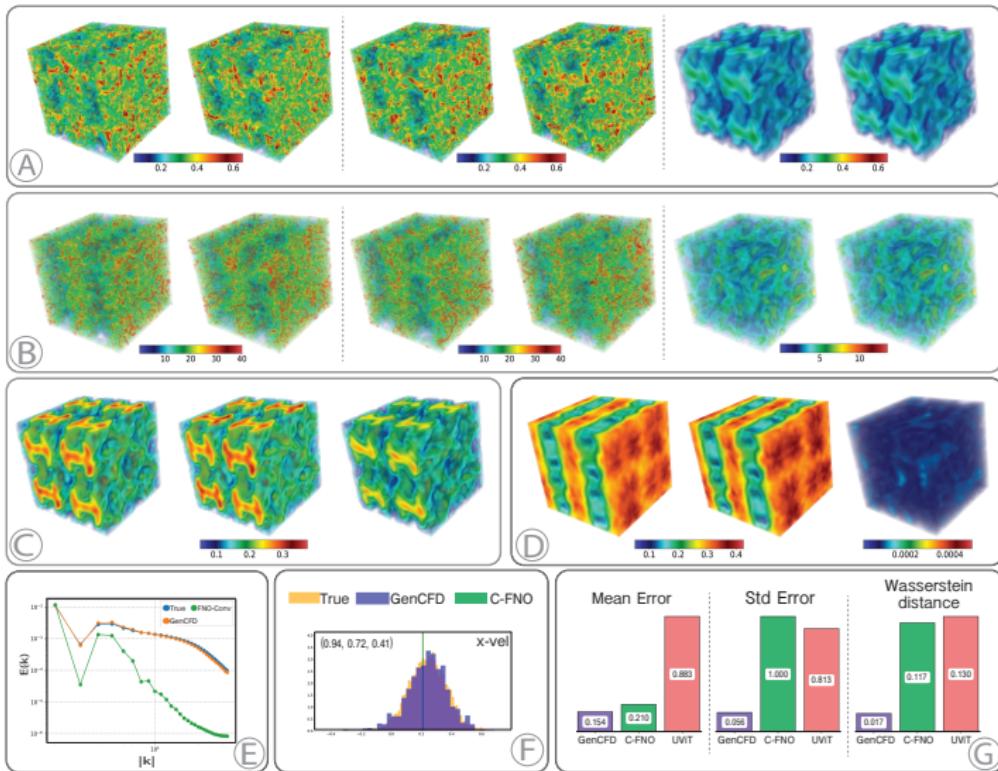
$$du_\tau = 2 \left(\frac{\dot{\sigma}_\tau}{\sigma_\tau} + \frac{\dot{s}_\tau}{s_\tau} \right) d\tau - 2s_\tau \frac{\dot{\sigma}_\tau}{\sigma_\tau} D_\theta(\Delta t, u_{\tau+1}, \bar{u}, \sigma_\tau) d\tau + s\sqrt{2\dot{\sigma}_\tau} d\hat{W}_\tau$$

- **Denoiser** minimizes $\mathbb{E}\|u(t_n, \bar{u}) - D_\theta(u(t_n, \bar{u}) + \eta, \bar{u}, \sigma)\|$

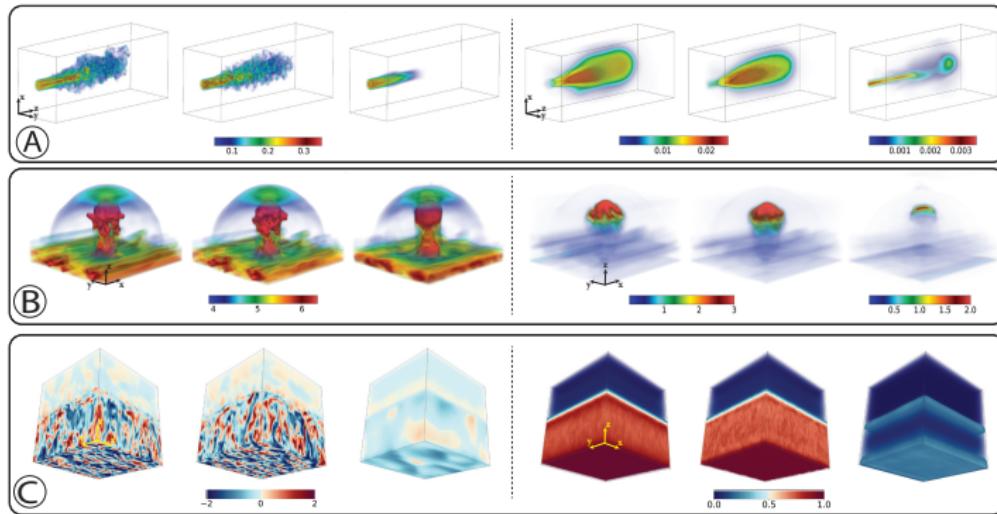


- GenCFD provably approximates the Conditional Distribution !!

Taylor-Green Results

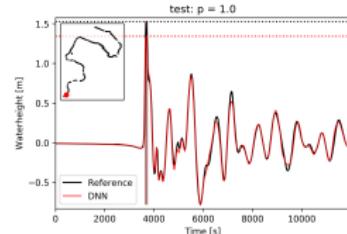
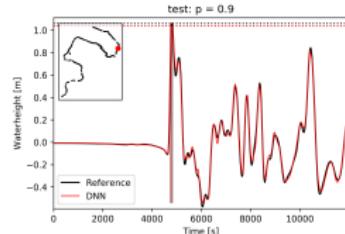
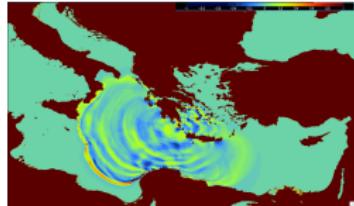


GenCFD works very well for Realworld Flows



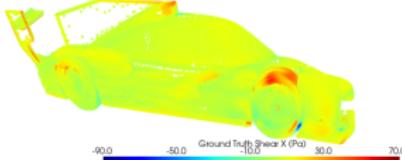
- ▶ Nozzle Jet: 3.5 hrs (LBM) vs. **GenCFD**: 1.45s
- ▶ Cloud-Shock: 5 hrs (FVM) vs. **GenCFD**: 0.45s
- ▶ Conv. Boundary Layer: 13.3 hrs (FDM) vs. **GenCFD**: 3.8s

What about the Use Cases ?

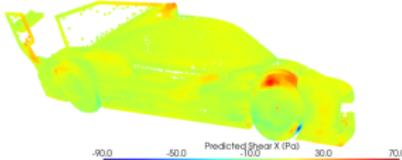


- ▶ AI Tsunami Simulation takes 10^{-3} secs (vs. 1 hr)
- ▶ AI RaceCar Simulation takes 10^{-2} secs (vs. 500 Node hrs)

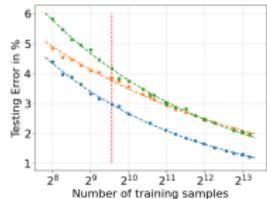
Mesh with Ground Truth Shear (X)



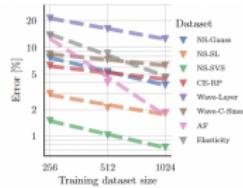
Mesh with Predicted Shear (X)



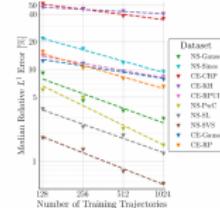
Where's the CAVEAT ?



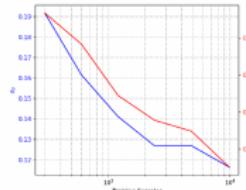
CNO/FNO



RIGNO



GAOT

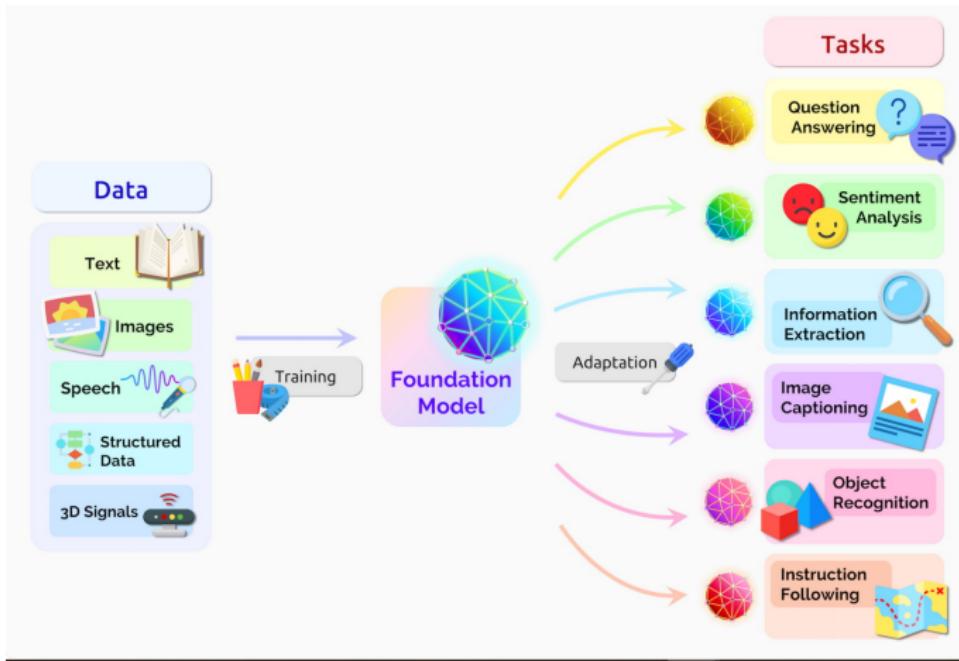


GenCFD

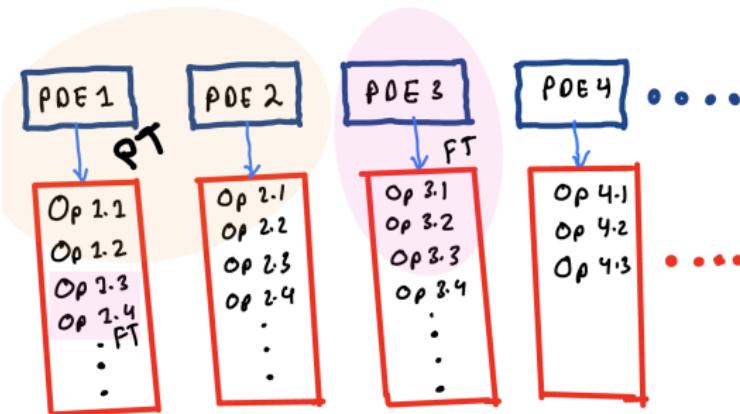
- ▶ Models **Scale** with sample size: $\mathcal{E} \sim N^{-\alpha}$ but with α **small**
- ▶ Even more pessimistic rates with theory ¹
- ▶ ML models require **Big Data**: $\mathcal{O}(10^3) - \mathcal{O}(10^4)$ training samples per Task
- ▶ Very Difficult to obtain Data for PDEs.
- ▶ How to make models **much more Sample Efficient** ?

¹ Lanthaler, SM, Karniadakis, 2022, De Ryck, SM, 2023

Foundation Models are the Key for Text/Vision!!

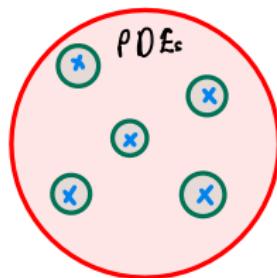


What would a Foundation Model for PDEs look like ?

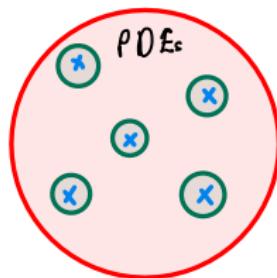


- ▶ **Op**: Operators need PDE + Data.
- ▶ **PT**: Pretraining.
- ▶ **FT**: Finetuning (Adaptation)

Can it Work ?

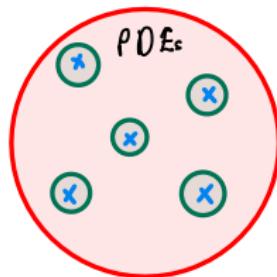


Can it Work ?



- Lets try it out !!

Can it Work ?

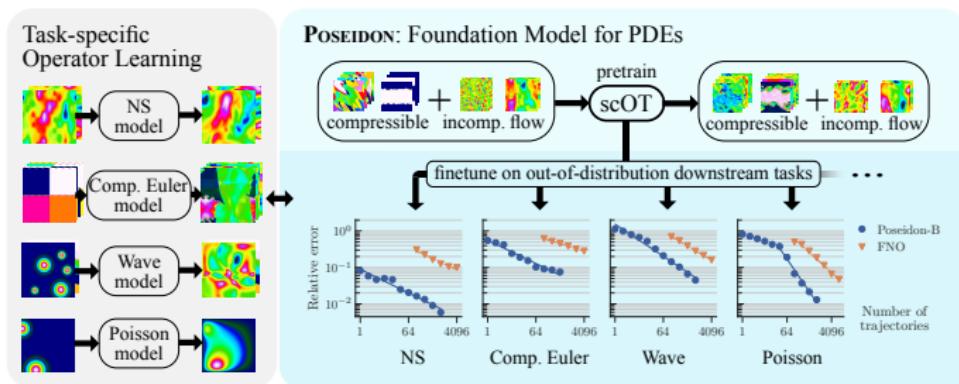


- Lets try it out !!

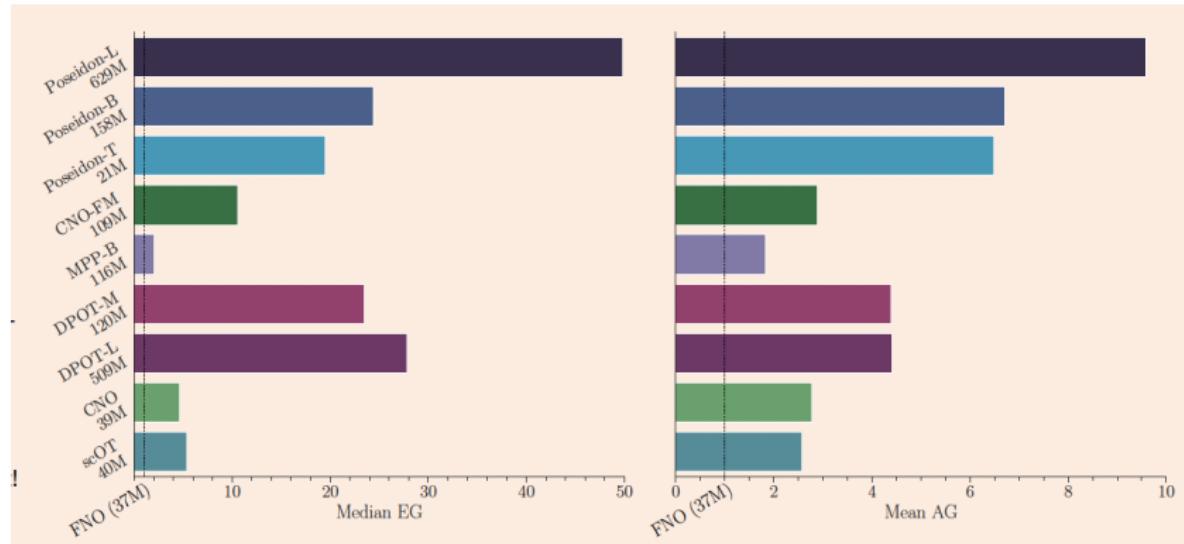


POSEIDON

- ▶ PDE foundation model of Herde, SM et. al, 2024
- ▶ Pretrained on Euler + Navier-Stokes Eqns in 2-D.
- ▶ Finetuned on 15 Unseen Tasks
- ▶ Paper + Code: <https://github.com/camlab-ethz/poseidon>
- ▶ Model + Datasets: <https://huggingface.co/camlab-ethz>

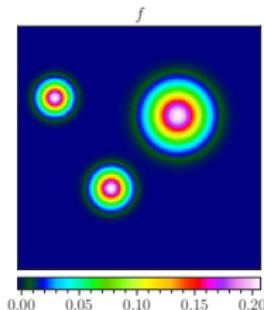


Summary of Performance on all Downstream Tasks

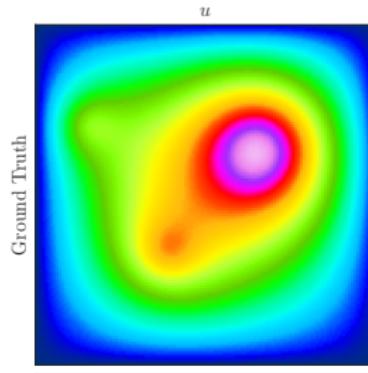


Deep Dive: Poisson Eqn.

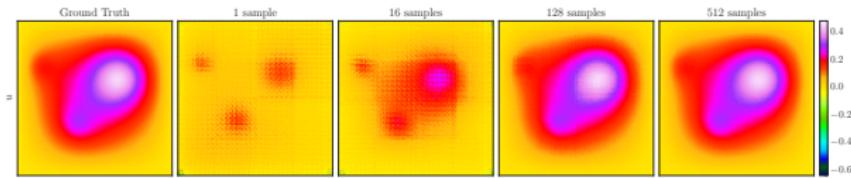
- Input (Source):



- Output (Solution):

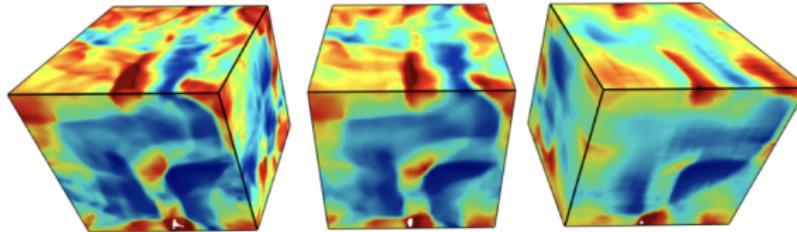


How much Physics has been learnt in PreTraining ?



Ongoing: Scaling up POSEiDON

- ▶ Increase **Model Size** by 4x: 2.5B Model.
- ▶ Increase **Dataset Size** by 10 – 50 x
- ▶ Augment Model **Features**:
 - ▶ 3D
 - ▶ Unstructured (point cloud) inputs
 - ▶ Genuine Multiphysics Training
 - ▶ Diffusion model for multiscale problems
 - ▶ PDE symbolic information



Summary+ Outlook

- ▶ ML/AI model can be potential **Neural PDE Solvers** (PINNs)
- ▶ Training is intrinsically **Ill-conditioned**.
- ▶ ML/AI models are effective **Neural PDE Surrogates**:
 - ▶ **Neural Operators** (CNO, scOT) for PDEs on Cartesian Domains.
 - ▶ **Graphs+Transformers** (GAOT) for Arbitrary Domains.
 - ▶ **Diffusion Models** (GenCFD) for Multiscale, Chaotic PDEs.
- ▶ **Sample Efficiency** is the main challenge.
- ▶ **Foundation Models** (Poseidon) can address it.
- ▶ They need to be Scaled Up significantly.